Prashant Pandey

Research Interest

My research interests lie at the intersection of Systems and Algorithms. I design and build theoretically well-founded data structures for big data problems in computational biology, streaming, and storage.

Contact Information

900 S Negley Avenue, Apt #6
Pittsburgh, PA - 15232
Website `https://prashantpandey.github.io`
Google Scholar `https://goo.gl/Fz82hB`

E-mail: ppandey2@cs.cmu.edu
Phone: (+1) 631-949-6948

Education

Stony Brook University, Stony Brook, NY — December 2018
PhD, Computer Science — *GPA (3.8/4.0)*
Thesis: Fast and Space-Efficient Maps: Shrinking Big Data Down to Size
Advisors: Prof. Michael Bender and Prof. Rob Johnson

University of Pune, Pune, India — August 2007 - June 2011
Bachelor of Engineering, Information Technology — *First class with distinction*
Ranked 1st in college and 7th across University

Work Experience

Carnegie Mellon University, Pittsburgh, PA — December 2018 - Present
Postdoctoral Scholar, School of Computer Science
Advisor: Prof. Carl Kingsford

Stony Brook University, Stony Brook, NY — August 2014 - October 2018
Research Assistant, Applied Algorithms Lab

TIBCO Inc., Pune, India — July 2011 - June 2013
Software Developer, Cloud Platform

Internships

Google, Manhattan, NY — May 2017 - August 2017
Research Intern, Google Spanner

Google, Kirkland, WA — May 2016 - August 2016
Research Intern, Google Could Infrastructure

Intel Labs, Portland, OR — May 2015 - August 2015
Research Intern, Security and Privacy Lab

Intel Labs, Portland, OR — May 2014 - August 2014
Research Intern, Security and Privacy Lab

Awards and Achievements

- Catacosinos Fellowship — 2018
- Best Paper Award FAST 2016 — 2016
- Runner's Up to Best Paper FAST 2015 — 2015
- A Special CS Department Chair Fellowship, Stony Brook University — 2013
- University Rank Holder, University of Pune — 2011
- Academic Excellence Scholarship, University of Pune. — 2009, 2010, 2011

Press Articles on Research

Finding a Needle in a Field of Haystacks. Cell Systems publishes research on Mantis, a new sequencing search tool. *July 2018*
Link: `https://goo.gl/LJopwR`

Our computational biology research got mentioned on VMware Research blog. *April 2018*
Link: `https://blogs.vmware.com/research/2018/04/18/scaling-computational-biology-vmware/`

A general purpose counting filter: making every bit count. The Morning Paper. *August 2017*
Link: `https://goo.gl/nReGcF`

Papers under submission

Timely Reporting of Heavy Hitters using External Memory
Prashant Pandey, Michael A. Bender, Jonathan W. Berry, Martin Farach-Colton, Rob Johnson, Thomas M. Kroeger, Cynthia A. Phillips, Shikha Singh

Publications

Locality Sensitive Hashing for the Edit Distance *ISMB 2019*
Guillaume Marais, Dan DeBlasio, **Prashant Pandey**, and Carl Kingsford

The Dictionary Problem, Optimal Searching, and Asymptotic Distortions of the DAM *SPAA 2019*
Michael A. Bender, Alex Conway, Martin Farach-Colton, William Jannen, Yizheng Jiao, Rob Johnson, Eric Knorr, Sara McAllister, Nirjhar Mukherjee, **Prashant Pandey**, Donald E. Porter, Jun Yuan, Yang Zhan

The Online Event-Detection Problem *arXiv 2019*
Michael A. Bender, Jonathan W. Berry, Martin Farach-Colton, Rob Johnson, Thomas M. Kroeger, **Prashant Pandey**, Cynthia A. Phillips, Shikha Singh

An Efficient, Scalable and Exact Representation of High-Dimensional Color Information Enabled via de Bruijn Graph Search Problem *RECOMB 2019*
Fatemeh Almodaresi, **Prashant Pandey**, Michael Ferdman, Rob Johnson, Rob Patro

Buffered Count-Min Sketch on SSD: Theory and Experiments *ESA 2018*
Mayank Goswami, Dzejla Medjedovic, Emina Mekic, **Prashant Pandey**

Mantis: A Fast, Small, and Exact Large-Scale Sequence-Search Index *RECOMB 2018 Cell Systems 2018*
Prashant Pandey, Fatemeh Almodaresi, Michael A. Bender, Michael Ferdman, Rob Johnson, and Rob Patro

Rainbowfish: A Succinct Colored de Bruijn Graph Representation *WABI 2017*
Fatemeh Almodaresi, **Prashant Pandey**, and Rob Patro

deBGR: An Efficient and Near-Exact Representation of the Weighted de Bruijn Graph *ISMB 2017 BIOINFORMATICS 2017*
Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

Squeakr: An Exact and Approximate k-mer Counting System *BIOINFORMATICS 2017*
Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

A General-Purpose Counting Filter: Making Every Bit Count *SIGMOD 2017*
Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

A Fast x86 Implementation of Select *arXiv 2017*
Prashant Pandey, Michael A. Bender, and Rob Johnson

Writes Wrought Right, and Other Adventures in File System Optimization *TOS 2016*
Jun Yuan, Yang Zhan, William Jannen, **Prashant Pandey**, Amogh Akshintala, Kanchan Chandnani, Pooja Deo, Zardosht Kasheff, Michael Bender, Martin Farach-Colton, Rob Johnson, Bradley C. Kuszmaul, and Donald E. Porter

Optimizing Every Operation in a Write-Optimized File System *FAST 2016*
Jun Yuan, Yang Zhan, William Jannen, **Prashant Pandey**, Amogh Akshintala, Kanchan Chandnani, Pooja Deo, Zardosht Kasheff, Michael Bender, Martin Farach-Colton, Rob Johnson, Bradley C. Kuszmaul, and Donald E. Porter [**Best Paper Award**]

BetrFS: Write-Optimization in a Kernel File System *TOS 2015*
William Jannen, Jun Yuan, Yang Zhan, Amogh Akshintala, John Esmet, Yizheng Jiao, Ankur Mittal, **Prashant Pandey**, Phaneendra Reddy, Leif Walsh, Michael A. Bender, Martin Farach-Colton, Rob Johnson, Bradley C. Kuszmaul, and Donald E. Porter

BϵtrFS: A Right-Optimized Write-Optimized File System *FAST 2015*
William Jannen, Jun Yuan, Yang Zhan, Amogh Akshintala, John Esmet, Yizheng Jiao, Ankur Mittal, **Prashant Pandey**, Phaneendra Reddy, Leif Walsh, Michael A. Bender, Martin Farach-Colton, Rob Johnson, Bradley C. Kuszmaul, and Donald E. Porter [**Runner up to Best Paper**]

Patent Applications

Instructions that Facilitate the Implementation of the Fork System Call in Processes using Software Guard Extensions *March 2015*
Prashant Pandey, Mona Vij, Somnath Chakrabarti, Krystof C. Zmudzinski

Apparatus and Method For Implementing a Forked System Call in a System with a Protected Region *March 2015*
Prashant Pandey, Mona Vij, Somnath Chakrabarti, Krystof C. Zmudzinski

Invited Talks

Compact Representation of Annotated de Bruijn Graphs
Berkeley Lab, Berkeley CA, January 2018

deBGR: An Efficient and Near-Exact Representation of the Weighted de Bruijn Graph
Google Research, NY, September 2017
VMWare Research, Palo Alto CA, Aug 2017

Intel Software Guard Extensions (SGX)
Sandia National Laboratories, Livermore CA, August 2015

Conference Talks

Timely Reporting of Heavy Hitters using External Memory
Theoretical Foundations of Storage Systems 2019, Dagstuhl, Germany

Buffered Count-Min Sketch on SSD: Theory and Experiments
ESA 2018, Helsinki, Finland

Mantis: A Fast, Small, and Exact Large-Scale Sequence-Search Index
RECOMB 2018, Paris, France

Scheduling Problems in Write-Optimized Key-Value Stores
New Challenges in Scheduling Theory 20018, Aussois, France

deBGR: An Efficient and Near-Exact Representation of the Weighted de Bruijn Graph
ISMB 2017, Prague, Czech Republic

A General-Purpose Counting Filter: Making Every Bit Count
SIGMOD 2017, Chicago, IL

Professional Service

- Program Committee — *ESA 2019*
- Subreviewer — *CIAC 2019*
- Reviewer — *IEEE Access*
- Reviewer — *Bioinformatics*
- Reviewer — *Transactions of Databases*

Teaching Experience

Teaching Assistant, CS Dept, Stony Brook University

- CSE 548: Analysis of Algorithms — *Fall 2015*
- CSE 535: Asynchronous Systems — *Fall 2015*
- CSE 110: Introduction to Computer Science (Advanced Java) — *Spring 2014*
- CSE 110: Introduction to Computer Science (Advanced Java) — *Fall 2013*