


# Vector Search for Large-Scale Genomic Discovery

VecDB@ICML2025




Prashant Pandey, Northeastern University, Boston USA  
<https://prashantpandey.github.io/>



# A typical genomic pipeline



# Assembled data is hugely lossy





# Assembled data is hugely lossy



Data Reduction: 500 GB → 1 GB (500x compression)

# Raw sequencing data contains biological diversity information

---



## Missing Species

We only have assembled genomes for a tiny fraction of Earth's ~8.7 million species. Most can't even be grown in labs!

## Population Diversity

Reference genomes reduce populations to single sequences, losing massive amounts of genetic variation and expression data.

A lot of variability information is lost during assembly. And a lot of raw sequencing data never gets assembled.



\*Mora, Camilo, et al. "How many species are there on Earth and in the ocean?." PLoS biology <https://doi.org/10.1371/journal.pbio.1001127>

# Raw sequencing data can unlock biological insights

Q. What if I find a new putative disease-related transcript, and want to see if it appeared in other biological samples?

Q. What if I discover a new fusion event in a particular cancer subtype and want to know if it is common among samples with this subtype?

Q. What if I find an unexpected bacterial contaminant in my data; which other samples might contain this?



```
> Analyzing sequences...
> ATCGATCGATCG
> GCTAGCTAGCTA
> Pattern detected!
> Variant found: SNP
> Gene expression: UP
> Disease risk: LOW
> Phylogeny: confirmed
> Analysis complete!
> Generating report...
```

"What biological  
insights can I  
discover?"



# Raw sequencing data can unlock biological insights

Q. What if I find a new putative disease-related transcript, and want to see if it appeared in other biological samples?


Q. What if I discover a new fusion event in a particular cancer subtype and want to know if it is common among samples with this subtype?

Q. What if I find an unexpected bacterial contaminant in my data; which other samples might contain this?

A. I need to perform string searches through tons of raw sequencing data.



# Sequence Read Archive (SRA) is growing rapidly



SRA is a publicly available dataset from NIH containing **raw sequencing data**

# Basic Local Alignment Search Tool (BLAST)

---

# Basic Local Alignment Search Tool (BLAST)

Contrast this situation with searching assembled, curated genomes, for which we have an excellent tool; BLAST

The screenshot shows the NCBI BLASTN search interface. On the left, the 'Enter Query Sequence' panel contains a text input field with the sequence: TGAAAAGGGTAACCTCAAAGCTAAAAGCCAAAGAAGGGGAAGCCCCATTGCAGCCGCAAC CCTGTCTTGTCAAGAGGAATTGGCAGGTATCCCGATC. Below it, there are options to upload a file or enter a job title. A large blue 'BLAST' button is centered at the bottom of the input area. To the right, the 'Sequences producing significant alignments' table is displayed, showing 10 entries. Each entry includes a checkbox, the sequence name, its length, identity percentage, E-value, and accession number.

	Max score	Total score	Query cover	E value	Ident	Accession
<input type="checkbox"/> Eukaryotic synthetic construct chromosome 18	185	371	100%	2e-43	100.00%	<a href="#">CP034496.1</a>
<input type="checkbox"/> PREDICTED: Pan paniscus 60S ribosomal protein L6-like (LOC100976413), mRNA	185	185	100%	2e-43	100.00%	<a href="#">XM_008963989.2</a>
<input type="checkbox"/> PREDICTED: Pan paniscus 60S ribosomal protein L6 pseudogene (LOC100995849), misc_RNA	185	185	100%	2e-43	100.00%	<a href="#">XR_610957.3</a>
<input type="checkbox"/> PREDICTED: Pan paniscus 60S ribosomal protein L6 (LOC100995836), mRNA	185	185	100%	2e-43	100.00%	<a href="#">XM_003812574.3</a>
<input type="checkbox"/> PREDICTED: Pan troglodytes 60S ribosomal protein L6 pseudogene (LOC737972), misc_RNA	185	185	100%	2e-43	100.00%	<a href="#">XR_680356.3</a>
<input type="checkbox"/> PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X8, mRNA	185	185	100%	2e-43	100.00%	<a href="#">XM_024347583.1</a>
<input type="checkbox"/> PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X7, mRNA	185	185	100%	2e-43	100.00%	<a href="#">XM_024347582.1</a>
<input type="checkbox"/> Human ORFeome Gateway entry vector pENTR223-RPL6, complete sequence	185	185	100%	2e-43	100.00%	<a href="#">LT737273.1</a>
<input type="checkbox"/> PREDICTED: Gorilla gorilla gorilla ribosomal protein L6 (RPL6), transcript variant X5, mRNA	185	185	100%	2e-43	100.00%	<a href="#">XM_019038370.1</a>

# Basic Local Alignment Search Tool (BLAST)

Contrast this situation with searching assembled, curated genomes, for which we have an excellent tool; BLAST

The screenshot shows the NCBI BLAST search interface. On the left, the "Enter Query Sequence" panel contains a text input field with the sequence: TGAAAAGGGTAACCTCAAAGCTAAAAGCCAAAGAAGGGGAAGCCCCATTGCAGCCGCAAC CCTGTCTTGTCAAGAGGAATTGGCAGGTATCCCGATC. Below it, there are options to "Or, upload file" and "Job Title". A "BLAST" button is centered below the input field. On the right, the "Sequences producing significant alignments" table displays search results. The table has columns for Max score, Total score, Query cover, E value, Ident, and Accession. The results list various entries, mostly predicted or eukaryotic synthetic constructs, with scores ranging from 185 to 371.

	Max score	Total score	Query cover	E value	Ident	Accession
<a href="#">Eukaryotic synthetic construct chromosome 18</a>	185	371	100%	2e-43	100.00%	<a href="#">CP034496.1</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6-like (LOC100976413), mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_008963989.2</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6 pseudogene (LOC100995849), misc_RNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XR_610957.3</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6 (LOC100995836), mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_003812574.3</a>
<a href="#">PREDICTED: Pan troglodytes 60S ribosomal protein L6 pseudogene (LOC737972), misc_RNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XR_680356.3</a>
<a href="#">PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X8, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_024347583.1</a>
<a href="#">PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X7, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_024347582.1</a>
<a href="#">Human ORFeome Gateway entry vector pENTR223-RPL6, complete sequence</a>	185	185	100%	2e-43	100.00%	<a href="#">LT73723.1</a>
<a href="#">PREDICTED: Gorilla gorilla gorilla ribosomal protein L6 (RPL6), transcript variant X5, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_019038370.1</a>

Essentially, the “Google of Genomics”

**Basic local alignment search tool**  
SF Altschul, W Gish, W Miller, EW Myers... - Journal of molecular ... , 1990 - Elsevier  
... A new approach to rapid **sequence** comparison, **basic local alignment search tool** (BLAST), directly approximates **alignments** that optimize a measure of **local** similarity, the maximal ...  
☆ Save ⚡ Cite Cited by 117758 Related articles All 41 versions

# Basic Local Alignment Search Tool (BLAST)

Contrast this situation with searching assembled, curated genomes, for which we have an excellent tool; BLAST

The screenshot shows the NCBI BLASTN search interface. On the left, the 'Enter Query Sequence' panel contains a text input field with the sequence: TGAAAAGGGTAACCTCAAAGCTAAAAGCCAAAGAAGGGGAAGCCCCATTGCAGCCGCAAC CCTGTCTTGTCAAGAGGAATTGGCAGGTATCCCGATC. Below it are options to 'Or, upload file' and 'Job Title'. A 'BLAST' button is centered below the input field. On the right, the 'Sequences producing significant alignments' panel displays a table of search results. The table has columns for Max score, Total score, Query cover, E value, Ident, and Accession. The results list various homologous sequences from different species, mostly predicted or synthetic constructs, with scores ranging from 185 to 371.

	Max score	Total score	Query cover	E value	Ident	Accession
<a href="#">Eukaryotic synthetic construct chromosome 18</a>	185	371	100%	2e-43	100.00%	<a href="#">CP034496.1</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6-like (LOC100976413), mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_008963989.2</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6 pseudogene (LOC100995849), misc_RNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XR_610957.3</a>
<a href="#">PREDICTED: Pan paniscus 60S ribosomal protein L6 (LOC100995836), mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_003812574.3</a>
<a href="#">PREDICTED: Pan troglodytes 60S ribosomal protein L6 pseudogene (LOC737972), misc_RNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XR_680356.3</a>
<a href="#">PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X8, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_024347583.1</a>
<a href="#">PREDICTED: Pan troglodytes ribosomal protein L6 (RPL6), transcript variant X7, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_024347582.1</a>
<a href="#">Human ORFeome Gateway entry vector pENTR223-RPL6, complete sequence</a>	185	185	100%	2e-43	100.00%	<a href="#">LT73723.1</a>
<a href="#">PREDICTED: Gorilla gorilla gorilla ribosomal protein L6 (RPL6), transcript variant X5, mRNA</a>	185	185	100%	2e-43	100.00%	<a href="#">XM_019038370.1</a>

Essentially, the “Google of Genomics”

**Basic local alignment search tool**  
SF Altschul, W Gish, W Miller, EW Myers... - Journal of molecular ... , 1990 - Elsevier  
... A new approach to rapid **sequence** comparison, **basic local alignment search tool** (BLAST), directly approximates **alignments** that optimize a measure of **local** similarity, the maximal ...  
☆ Save ⚡ Cite Cited by 117758 Related articles All 41 versions

BLAST performs substring matching (based on seed-search-align) using traditional succinct string data structures.

# Why can't we use BLAST for searching “raw” data?

---



## Fragmented Patterns

The sequence you're looking for might be spread across multiple reads, making it impossible for BLAST to find as a contiguous match.

## Scale Limitations

BLAST algorithms and data structures simply don't scale to handle millions of raw sequencing experiments efficiently.

# Reframing the problem as vector search

Solomon and Kingsford 2016 reframed the problem, and suggested a direction...



## Fast search of thousands of short-read sequencing experiments

Brad Solomon & Carl Kingsford

*Nature Biotechnology* **34**, 300–302 (2016)

doi:10.1038/nbt.3442

[Download Citation](#)

Received: 28 April 2015

Accepted: 23 November 2015

Published online: 08 February 2016


### Proposal:

Represent each sample from SRA as a set of *tokens*

Similarity of *token* composition — similar sequence — small edit distance

Returns “yes/no” results for individual samples — “yes” results can be searched using traditional methods

# $K$ -mers as search primitives



- For a given molecule (string), a  $k$ -mer is simply a  $k$ -length sub-string
- Akin to n-grams used in NLP (except DNA/RNA have no natural “tokens”)
- **Idea:** Similarity of  $k$ -mer composition — similar sequence

# Sample discovery problem

---

# Sample discovery problem


## Query transcript

🔍 ATCGATCGATCGAATCG


*k*-mers (*k*=5)

ATCGA	TCGAT	CGATC
GATCG	ATCGA	TCGAA
CGAAT	GAATC	AATCG


# Sample discovery problem



# Sample discovery problem




# Sample discovery problem




Return all samples that contain at least some  $\theta$  fraction of  $k$ -mers present in the query transcript


# Sample discovery → Vector search



$K$ -mer representation of sequences is employed as vector embeddings.  $k$  ranges between 21-31



Hamming distance between  $k$ -mer embeddings is employed as a proxy for sequence similarity



Distance threshold is employed as a proxy for the top- $k$  nearest neighbor


Underlying biological generative model claims that similar species have similar  $k$ -mer content.

Hamming distance is often not a good proxy for sequence similarity\*. (More on this later in the talk.)

\*Guillaume Marcais, Dan DeBlasio, Prashant Pandey, Carl Kingsford "Locality Sensitive Hashing for the Edit Distance" ISMB 2019

# Several indexing methods for raw sequence search

## Color aggregative methods



## $k$ -mer aggregative methods



**K-mer set  
data structure**

Hash tables, CQF, BWT, BF trie

Bloom filters

**Aggregation  
data structure**

Color matrices

Search tree/forests, Bloom filter  
matrices

**Method  
names**

Mantis, SeqOthello, Bifrost,  
Metanot, BFT, VARI,


SBT (variants), BIGSI, COBS,  
RAMBO

Vibrant area with exciting work over the past several years; excellent review by Marchet et al. 2021.


# $K$ -mer aggregative methods: SBT

[Solomon & Kingsford 2016]

- A binary-tree of Bloom filters, where leaves represent the  $k$ -mer set of a single sample
- Bloom filter of parent is logical union (= bitwise OR) of children
- Check both children, stop descending into tree when  $\theta$  threshold is not satisfied




# Several limitations in Bloom filter-based methods



All Bloom filters are of same size — creating skewness in the query results



Bloom filters cause false positives in the results



Hard to associate any attributes such as abundance, position etc.

SRA sample sizes range between few MBs to a few GBs. Results in high variance in Bloom filter false-positive rates.

# A fundamentally different approach

---

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”

# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

K-mer index

## Squeakr: an exact and approximate k-mer counting system

Bioinformatics 2018

Prashant Pandey<sup>1,\*</sup>, Michael A. Bender<sup>1</sup>, Rob Johnson<sup>1,2</sup> and Rob Patro<sup>1</sup>

<sup>1</sup>Department of Computer Science, Stony Brook University, Stony Brook, NY 11790, USA and <sup>2</sup>VMware Research, Palo Alto, CA 94304, USA

# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

Interesting observation  
about patterns of k-mer occurrence



## Rainbowfish: A Succinct Colored de Bruijn Graph Representation\*

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>2</sup>, and Rob Patro<sup>3</sup>

WABI 2017

K-mer index

## Squeakr: an exact and approximate k-mer counting system

Bioinformatics 2018

Prashant Pandey<sup>1,\*</sup>, Michael A. Bender<sup>1</sup>, Rob Johnson<sup>1,2</sup> and Rob Patro<sup>1</sup>

<sup>1</sup>Department of Computer Science, Stony Brook University, Stony Brook, NY 11790, USA and <sup>2</sup>VMware Research, Palo Alto, CA 94304, USA

# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

Interesting observation  
about patterns of k-mer occurrence



## Rainbowfish: A Succinct Colored de Bruijn Graph Representation\*

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>2</sup>, and Rob Patro<sup>3</sup>

WABI 2017

“I bet we can exploit  
that for large-scale search”

K-mer index

## Squeakr: an exact and approximate k-mer counting system

Bioinformatics 2018

Prashant Pandey<sup>1,\*</sup>, Michael A. Bender<sup>1</sup>, Rob Johnson<sup>1,2</sup> and Rob Patro<sup>1</sup>

<sup>1</sup>Department of Computer Science, Stony Brook University, Stony Brook, NY 11790, USA and <sup>2</sup>VMware Research, Palo Alto, CA 94304, USA

## Mantis: A Fast, Small, and Exact Large-Scale Sequence-Search Index

Prashant Pandey<sup>1</sup>, Fatemeh Almodaresi<sup>1</sup>, Michael A. Bender<sup>1</sup>, Michael Ferdman<sup>1</sup>, Rob Johnson<sup>2,1</sup>, and Rob Patro<sup>1</sup>

RECOMB 2018 & Cell Systems

# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

Interesting observation  
about patterns of k-mer occurrence



## Rainbowfish: A Succinct Colored de Bruijn Graph Representation\*

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>2</sup>, and Rob Patro<sup>3</sup>

WABI 2017

“I bet we can exploit  
that for large-scale search”

K-mer index

## Squeakr: an exact and approximate k-mer counting system

Prashant Pandey<sup>1,\*</sup>, Michael A.

<sup>1</sup>Department of Computer Science, Stony Brook  
Palo Alto, CA 94304, USA

## An Efficient, Scalable and Exact Representation of High-Dimensional Color Information Enabled via de Bruijn Graph Search

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>1</sup>, Michael Ferdman<sup>1</sup>, Rob Johnson<sup>2,1</sup>, and Rob Patro<sup>1</sup>


RECOMB 2019 & JCB 2020

## Mantis: A Fast, Small, and Exact Large-Scale Sequence-Search Index

Prashant Pandey<sup>1</sup>, Fatemeh Almodaresi<sup>1</sup>, Michael A. Bender<sup>1</sup>, Michael Ferdman<sup>1</sup>, Rob Johnson<sup>2,1</sup>, and Rob Patro<sup>1</sup>

RECOMB 2018 & Cell Systems

“I bet we can make  
it even smaller”



# A fundamentally different approach

Our initial idea: “The Bloom Filter is limiting. What can we get by replacing it with a *better* filter ?”



## A General-Purpose Counting Filter: Making Every Bit Count

Prashant Pandey, Michael A. Bender, Rob Johnson, and Rob Patro

SIGMOD 2017

Interesting observation  
about patterns of k-mer occurrence



## Rainbowfish: A Succinct Colored de Bruijn Graph Representation\*

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>2</sup>, and Rob Patro<sup>3</sup>

WABI 2017

“I bet we can exploit  
that for large-scale search”

K-mer index

## Squeakr: an exact and approximate k-mer counting system

Prashant Pandey<sup>1,\*</sup>, Michael A.

<sup>1</sup>Department of Computer Science, Stony Brook  
Palo Alto, CA 94304, USA

## An Efficient, Scalable and Exact Representation of High-Dimensional Color Information Enabled via de Bruijn Graph Search

Fatemeh Almodaresi<sup>1</sup>, Prashant Pandey<sup>1</sup>, Michael Ferdman<sup>1</sup>, Rob Johnson<sup>2,1</sup>, and Rob Patro<sup>1</sup>

RECOMB 2019 & JCB 2020

RECOMB 2018 & Cell Systems

## Mantis: A Fast, Small, and Exact Large-Scale Sequence-Search Index

Prashant Pandey<sup>1</sup>, Fatemeh Almodaresi<sup>1</sup>, Michael A. Bender<sup>1</sup>, Michael Ferdman<sup>1</sup>, Rob Johnson<sup>2,1</sup>, and Rob Patro<sup>1</sup>

“I bet we can make  
it even smaller”


Bioinformatics 2022

“I bet we can make  
it scale and updatable”

# Color aggregative methods: Mantis

[Pandey et al. 2018]

- Build a counting quotient filter for each input sample (can be different sizes based on the number of  $k$ -mers)
- CQF: **key**= $k$ -mer **value**=color class ID
- Combine them via multi-way merge
- Estimate a good ordering of color class IDs from first few million  $k$ -mers



\*Raman, et al. (2002). Succinct indexable dictionaries with applications to encoding k-ary trees and multisets. SODA

# Mantis is faster, smaller, and accurate than SSBT

Indexed 2,652 human BBB RNA-seq (gene expression) samples ~**4.5TB** of (Gzip compressed) data

**Table 1. Time and Space Measurement for Mantis and SSBT**

Tool	Mantis	SSBT
<b>Build time</b>	03 hr 56 min	97 hr
<b>Representation size.</b>	32 GB	39.7 GB

Query includes index loading

$\theta$  threshold for SSBT query


	Mantis	SSBT (0.7)	SSBT (0.8)	SSBT (0.9)
10 Transcripts	25 s	3 min 8 s	2 min 25 s	2 min 7 s
100 Transcripts	28 s	14 min 55 s	10 min 56 s	7 min 57 s
1000 Transcripts	1 min 3 s	2 hr 22 min	1 hr 54 min	1 hr 20 min

Mantis can be constructed ~24x faster than a comparable SSBT\* [Solomon & Kingsford 2017]

Mantis is ~6 – 109x faster than (in memory) SSBT

The final Mantis representation is ~20% smaller than the comparable SSBT representation.

# Where we are now




# Some remaining challenges

We demonstrate indexing on the order of  $10^3$  samples, we really want to index on the order of  $10^6$  samples

## Key observations:

- $K$ -mers grow at worst linearly
- Color classes increase super-linearly




Need a fundamentally better color class encoding; exploit coherence between rows of the color class matrix

# Consider the following color class graph

Each color class is a vertex in the graph


Every pair of color classes is connected by an edge whose weight is the hamming distance between the color class vectors



# Consider the following color class graph

Each color class is a vertex in the graph


Every pair of color classes is connected by an edge whose weight is the hamming distance between the color class vectors



# Consider the following color class graph

Each color class is a vertex in the graph

Every pair of color classes is connected by an edge whose weight is the hamming distance between the color class vectors




Unfortunately:

- 1) There are many color classes (> 1 Billion, full graph too big)
- 2) They are high-dimensional (# of samples), neighbor search is hard (LSH scheme seem to work poorly)

# De Bruijn graph (dBG) over $k$ -mers

Read 1: ....CAAAAT....

Read 2: ....CAAAAC....





Use the de Bruin graph (dBG) as an **efficient guide for near-neighbor search** in the space of color classes!  
dBG common in genomics. Nodes  $u, v$  are  $k$ -mers & are adjacent if  $k-1$  suffix of  $u$  is the  $k-1$  prefix of  $v$

# Mantis implicitly represents a colored dBG

Each CQF key represents a  $k$ -mer → can explicitly query neighbors

Each  $k$ -mer associated with color class id → vector of occurrences






$$\begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 0 & 0 \end{bmatrix}$$



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 0 \end{bmatrix}$$







$$\begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \end{bmatrix}$$



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 0 & 1 & 1 \\ 0 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$































































Complete CCG




Optimal MST




Complete CCG




Optimal MST




CCG derived from dbG



MST on our Graph




# The MST efficiently encodes related color classes




# The MST efficiently encodes related color classes

Augment with all 0 color class to guarantee one, connected MST




# The MST efficiently encodes related color classes

Augment with all 0 color class to guarantee one, connected MST




# The MST efficiently encodes related color classes

Augment with all 0 color class to guarantee one, connected MST



To reconstruct a vector, walk from your node to the root, flipping the parity of the positions you encounter on each edge.

# The MST approach scales very well



MST-based  
compression of bit  
vectors

Dataset	# samples	MST					$\frac{\text{size}(MST)}{\text{size}(RRR)}$
		RRR matrix	Total space	Parent vector	Delta vector	Boundary bit-vector	
dataset from SBT / SSBT / Mantis paper	200	0.42	0.15	0.08	0.06	0.01	0.37
	500	1.89	0.46	0.2	0.24	0.03	0.24
	1,000	5.14	1.03	0.37	0.6	0.06	0.2
	2,000	14.2	2.35	0.71	1.5	0.14	0.17
	5,000	59.89	7.21	1.72	5.1	0.39	0.12
	10,000	190.89	16.28	3.37	12.06	0.86	0.085
	Blood, Brain, Breast (BBB)	15.8	2.66	0.63	1.88	0.16	0.17

Improvement  
over RRR improves  
with # of samples

# How does MST approach affect query time?

---

One concern is that replacing  $O(1)$  lookup with  
MST-based decoding will make lookup slow; does it?


# How does MST approach affect query time?

One concern is that replacing  $O(1)$  lookup with  
MST-based decoding will make lookup slow; does it?

Turns out a caching strategy (an LRU over popular internal nodes)  
keeps it just as fast as lookup in the RRR matrix

	Mantis with MST			Mantis		
	index load + query	query	space	index load + query	query	space
10 Transcripts	1 min 10 sec	0.3 sec	118GB	32 min 59 sec	0.5 sec	290GB
100 Transcripts	1 min 17 sec	8 sec	119GB	34 min 33 sec	11 sec	290GB
1000 Transcripts	2 min 29 sec	79 sec	120GB	46 min 4 sec	80 sec	290GB

# Where we are now



# Some remaining challenges

---


- We can scale to even larger datasets by compressing color class representation
- We demonstrate indexing on the order of  $10^3$  experiments, we really want to index on the order of  $10^5 - 10^6$
- We need to scale out of RAM and also support adding new experiments

## ***Key Observation:***

- We can take a static representation and make it updatable using the Bentley-Saxe construction [Bentley and Saxe (1980)]
- We can reduce the memory usage using ***minimizers***.

Need a ***fundamentally better*** construction which can support adding new samples and can scale out of RAM to disk.


# Mantis-LSM design



- Level 0 resides in RAM
- L<sub>1</sub>...L<sub>n</sub> remain on disk
- Levels grow in size exponentially

- **Minimizers** to partition the k-mer index on disk
- Helps to minimize RAM usage during merging and queries.

# Mantis-LSM performance



**Fig. 4.** Performance of the Dynamic Mantis update process. The spikes in time (Fig. a) and memory (Fig. b) happen when the cascading merge happens with deeper and thus larger indexes. Cumulative Time (Fig. c) shows the total time required to add all the samples up to the current one, and index size (Fig. d) is total size of the index

**LSM-Mantis can index up to 40K samples on a single machine**

# Mantis Distributed


- We now have a distributed Mantis (**publicly-hosted search service over SRA**)
- Each individual machine runs Mantis-LSM (written in RUST)
- Employ *minimizers* to efficiently distribute samples across machines
- Indexes **100K samples** from SRA (~200TB compressed data)
- **47 machines** each with 32GB RAM
- Support **real-time queries** (< 1 sec)

<https://query.bio/>


The screenshot shows the Mantis search interface. At the top, there's a logo of a mantis and the word "MANTIS". Below the logo, a search bar contains the query string "ACTTGAGTCGAGTTGACGCAAAAGTG". To the right of the search bar is a "Search" button and a "Choose FASTA" link. The main area is titled "Search Results:" and displays the following information:  
Query String: ACTTGAGTCGAGTTGACGCAAAAGTG  
Query ID: mg03138ihhk4rimb1gnqf  
Query Filter:  
Status: inprogress  
Waiting On: 4  
Initiated At: 7/18/2025, 12:35:28 AM  
Completed At: None  
Below this, there's a search bar with "Search..." and a "Download" button. At the bottom, there's a "Result" section showing "0-0 of 0 items" and navigation links for "1" and "of 1 page".

# Major insights from sample discovery problem

Domain specific knowledge enables to develop efficient solutions



K-mer embeddings  
appropriately capture  
sequence similarity



K-mer de Bruijn graph is  
a good proxy for the  
neighborhood graph

However, scale of the SRA is still the biggest challenge!

# $K$ -mer set representation ignores relative order

$x = \text{CCCCACCAACACACAAAAACCC} \rightarrow \begin{array}{l} \text{AAAA, AAAC, ACA, ACC, ACAA, ACAC,} \\ \text{ACCA, ACCC, CAAA, CAAC, CACA, CACC,} \\ \text{CCAA, CCAC, CCCA, CCCC} \end{array}$


$y = \text{AAAACACAACCCCCACCAAA} \rightarrow \begin{array}{l} \text{AAAA, AAAC, ACA, ACC, ACAA, ACAC,} \\ \text{ACCA, ACCC, CAAA, CAAC, CACA, CACC,} \\ \text{CCAA, CCAC, CCCA, CCCC} \end{array}$

$x$  and  $y$  have the same  $k$ -mer representation but low sequence similarity

$K$ -mer representation might not be the right proxy\*. Can introduce a large number of false positives.  
Need to evaluate the end-to-end accuracy of  $k$ -mer-based methods.

\*Guillaume Marcais, Dan DeBlasio, Prashant Pandey, Carl Kingsford "Locality Sensitive Hashing for the Edit Distance" ISMB 2019

# Employing ML/vector databases for genomic discovery



A ton of publicly-available data from SRA

# Major takeaways

---



- Sequence search over SRA promises to unlock deep biological insights
- Scalability is the biggest challenge; there's still a long way to go to index/search all of SRA
- Exploiting domain knowledge enables simpler, more scalable solutions
- Time to employ ML and vector databases to solve sequence search

# A special thanks to my collaborators

## Funding:




National  
Science  
Foundation



Jamshed Khan  
(UMD)



Fatemeh Almodaresi  
(OICR)



Rob Patro  
(UMD)



Mike Ferdman  
(Stony Brook)



Rob Johnson  
(VMware Research)



Michael Bender  
(Stony Brook)



<https://prashantpandey.github.io/>

